

A Reproducible Simulation Framework for Comparing Legislative Mechanism Bundles

Stress Tests for Productivity, Revision Moderation, Public Support, and Risk

ANONYMOUS AUTHOR(S)

Legislatures need enough productivity to act, but productivity alone can also mean volatility, agenda capture, or weakly supported lawmaking. This paper presents a reproducible computational framework for comparing selected legislative mechanism bundles under shared synthetic assumptions. The simulator routes the same generated legislatures and bill streams through legislative-style processes that vary agenda access, proposal revision, voting thresholds, alternative selection, public review, lobbying safeguards, affected-group protection, package bargaining, and reversibility.

The contribution is methodological: a reusable experimental architecture for diagnosing how legislative collective decision mechanisms shift the relationship among throughput, moderating revision, support, risk, and procedural cost under controlled synthetic worlds. The framework is designed to make tradeoffs inspectable rather than to rank real institutions. It reports a diagnostic dashboard separating productivity, revision moderation, public support, low-public-support enactment, risk control, administrative feasibility, and generated public benefit per submitted proposal. A main comparison campaign demonstrates how these metrics separate across representative mechanism families, while scripted diagnostics generate seed, ablation, manipulation-stress, chamber-structure, and empirical boundary checks. The reported results are synthetic, generator-dependent design hypotheses.

CCS Concepts: • **Human-centered computing** → **Collaborative and social computing theory, concepts and paradigms**; • **Computing methodologies** → **Modeling and simulation**; • **Applied computing** → *Law*.

Additional Key Words and Phrases: legislative institutions, institutional design, computational simulation, agenda control, revision moderation, collective decision making, voting systems

ACM Reference Format:

Anonymous Author(s). 2026. A Reproducible Simulation Framework for Comparing Legislative Mechanism Bundles: Stress Tests for Productivity, Revision Moderation, Public Support, and Risk. In *Proceedings of ACM Collective Intelligence 2026 (CI 2026)*. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

The motivating question is whether systematic comparison can distinguish legislative structures that merely pass bills from those that combine productivity, revision, and representative responsiveness. A legislature can be unified without being democratic; authoritarian systems can appear cohesive because disagreement is suppressed. In this framework, representative responsiveness is operationalized through generated public support, generated public benefit, revision behavior, and burden diagnostics for affected groups.

The central modeling choice is to treat final passage thresholds as only one part of a larger institutional process that includes proposal access, agenda control, revision, review, and correction. Burden-shifting passage rules are included only as stress tests because they expose how proposal access, objection rights, information, and correction interact with final voting thresholds.

This paper treats legislative institutions as collective intelligence architectures: procedural rules determine who can contribute proposals, which signals are amplified, how expertise and public support enter decisions, how disagreement

2026. Manuscript submitted to ACM

is resolved, and how errors are corrected. The simulator is therefore not a forecast of Congress but a tool for comparing institutional aggregation mechanisms under shared synthetic conditions.

The paper contributes four linked artifacts: a modular simulator for legislative mechanism bundles; a diagnostic dashboard separating productivity, support, risk, revision, and cost; a reproducible comparison campaign over selected mechanism families; and stress tests showing how generator assumptions, vote-kernel assumptions, and bounded manipulation probes affect the displayed tradeoffs.

The paper asks three questions. First, can a modular simulator distinguish throughput-producing rules from rules that produce moderated, broadly supported revisions under shared synthetic assumptions? Second, which mechanism families remain relevant when productivity, revision moderation, public support, risk, and administrative burden are evaluated separately? Third, which findings survive sensitivity checks on assumptions, mechanism ablations, and manipulation stressors?

This paper should be read as a framework and stress-test study, not as a validated model of Congress or a recommendation for constitutional reform. The reported results compare institutional mechanisms under shared synthetic assumptions. Empirical held-out and flow smoke tests screen for gross implausibility, but bill-topic opinion, affected-group evidence, bill-specific finance/lobbying influence, implementation outcomes, direct court/statute review, emergency court review, full statutory lineage, and cross-national institutional validation remain future work.

2 Related Work

The model draws on individual-level simulation, spatial voting, and institutional theory. Agent-based and computational political models are useful when heterogeneous actors interact under rules that create aggregate outcomes [de Marchi and Page 2014; Epstein 2006; Gilbert and Troitzsch 2005]. Spatial voting motivates one-dimensional ideal points and status-quo-relative utility [Downs 1957]. Legislative bargaining and agenda-setter models motivate explicit proposer identity and proposer gain [Baron and Ferejohn 1989; Romer and Rosenthal 1978]. Pivotal-politics and veto-player theories motivate comparing simple majority, supermajority, bicameral, and executive-veto systems [Krehbiel 1998; Tsebelis 2002]. Committee and agenda-control theory motivates treating floor denial, referral, information, and amendment control as outcomes rather than missing observations [Cox and McCubbins 2005; Gilligan and Krehbiel 1987; Krehbiel 1991; Shepsle 1979; Shepsle and Weingast 1987; Weingast and Marshall 1988].

Several tested mechanisms come from adjacent literatures. Lobbying is modeled as access, effort, information, public pressure, and participation bias rather than only direct vote buying [Austen-Smith 1993, 1995; de Figueiredo and Richter 2014; Gilens and Page 2014; Hall and Deardorff 2006; Hall and Wayman 1990]. Citizen-panel variants draw on deliberative polling and mini-public practice [Fishkin 2009; Organisation for Economic Co-operation and Development 2020]. Agenda-credit variants echo quadratic-voting and storable-vote work on expressing intensity under scarcity [Casella 2005; Lalley and Weyl 2018]. Reversibility variants draw on temporary legislation and sunset review [Gersen 2007; Ranchordas 2014]. Policy tournaments draw on Condorcet and agenda-manipulation problems [McKelvey 1976; Tideman 1987]. Model documentation follows ODD and ODD+D conventions [Grimm et al. 2010; Müller et al. 2013]. In collective intelligence terms, legislatures combine proposal generation, information filtering, preference aggregation, deliberative revision, and error correction. Existing legislative models often isolate one stage; this framework treats the stages as an interacting architecture. This paper contributes an executable comparison environment that treats institutional design as a modular collective intelligence architecture: the same synthetic actors and proposal streams are passed through different agenda, revision, voting, review, and correction mechanisms so diagnostic tradeoffs become inspectable.

3 Simulator

Each run creates a synthetic legislature, bill stream, scalar status quo, and institutional process. Legislators have ideal points $z_i \in [-1, 1]$, party labels, district preferences, and heterogeneous sensitivity to party pressure, constituency pressure, lobbying, reputation, and moderating revision. Bills have a proposer, policy position x_b , public support p_b , generated public benefit g_b , lobby pressure ℓ_b , salience, private gain, issue domain, public-benefit uncertainty, affected-group support, and concentrated harm. Enacted bills update one abstract status quo q_t , so votes evaluate policy movement rather than isolated proposals. Issue domains affect lobbying, harm, committees, and diagnostics, but the baseline model intentionally collapses policy movement into one dimension.

The voting kernel is status-quo-relative. A legislator's spatial utility is $u_i(x) = -(x - z_i)^2$. For bill b , the vote score is a weighted sum of ideology, party, public support, lobbying, revision moderation, and reputation terms:

$$s_i(b) = \sum_k w_k f_k(i, b, q_t) + \epsilon_i, \quad \Pr(Y_i = 1) = \sigma(\max(-12, \min(12, s_i(b)))).$$

The ideology term uses $u_i(x_b) - u_i(q_t)$, party pressure uses the same comparison at the party position, constituency pressure is proportional to $2(p_b - 0.5)$, lobbying is proportional to ℓ_b , and the revision-moderation term rewards moderation, public legitimacy, and incremental policy movement. The outcome is stochastic but seed-reproducible.

Institutional modules then transform the bill lifecycle:

$$(W_t, b, q_t) \rightarrow \text{access} \rightarrow \text{revision or selection} \rightarrow \text{vote} \rightarrow \text{review} \rightarrow (q_{t+1}, m_t).$$

Modules can deny access, route to committees, force discharge or open-calendar consideration, spend attention, apply lobbying safeguards, mediate text, select among alternatives, impose harm-weighted thresholds, run citizen or objection review, apply vetoes or judicial review, and register enacted laws for later correction. Post-enactment review tracks implementation delay, capacity, underfunding, nonenforcement, and renewal pressure as separate diagnostics.

The executable loop is: generate a world; replay its bill stream through each mechanism bundle; apply the access, revision/selection, vote, review, and status-quo-update modules for each bill; then aggregate diagnostics across worlds, cases, and seeds. This paired replay is what lets the framework compare procedural changes while holding generated actors and proposals fixed.

The stylized U.S.-like conventional benchmark uses House majority passage, a Senate 60 percent threshold, presidential veto, two-thirds override, and an agenda-access gate. The gate prevents the benchmark from treating every introduced bill as a floor bill; it schedules bills using observable proxies: public support, salience, policy stability, sponsor-party fit, partisan queue conflict, capture risk, concentrated harm, and uncertainty. It does not observe generated public benefit directly.

The baseline generator is moderation-friendly by construction: generated public benefit increases as $|x_b|$ decreases, before noise, private-gain penalties, and adversarial-case modifications. That assumption makes content-selection and moderating-revision mechanisms easier to reward. The results section therefore treats the ordering as a demonstration of framework behavior under the implemented generator, and the appendix reports alternative worlds where extreme proposals can be beneficial, moderation can dilute value, lobbying can supply information, public support can be biased, and majority support can coincide with concentrated harm.

Table 1. Key generator and vote-kernel assumptions visible in the main paper. Exact formulas, bounds, scenario constants, and weights are reported in the appendix and code. These are modeling choices, not empirical estimates.

Component	Core modeling assumption	Why it matters
Legislator ideal points	One-dimensional, faction-centered distribution with heterogeneous party, constituency, lobbying, reputation, and revision-moderation sensitivities.	Determines ideological conflict and makes moderation mostly movement along one policy line.
Status quo and bills	Bills are proposer-centered policy movements evaluated against a scalar status quo; enacted bills update that status quo.	Makes voting status-quo-relative, but collapses multidimensional politics into one abstract dimension.
Public benefit	Baseline g_b is moderation-linked: $E[g_b] \approx 0.34 + 0.34(1 - x_b) - 0.08 \text{ private}_b$, before noise and adversarial-world changes.	This is the most important generator vulnerability because it can advantage content-selection and moderating-revision mechanisms.
Public support	Support is synthetic and correlated with generated benefit, lobbying pressure, private gain, and noise.	Public-support diagnostics are not observed public opinion and inherit generator assumptions.
Lobby/private gain	Lobby pressure is partly tied to private gain, with alternative worlds where lobbying can provide useful information.	Drives capture diagnostics and makes anti-lobbying assumptions visible.
Harm/uncertainty	Concentrated harm and uncertainty rise with distance, salience, private gain, signal mismatch, lobbying, and noise.	Drives risk screens, harm-weighted thresholds, and low-support/risk-control metrics.
Vote kernel	Legislators combine status-quo-relative spatial utility, party pressure, public support, lobbying, a revision-moderation term, reputation, and seeded noise.	The revision-moderation score is a model proxy, not a neutral moral measure of good bargaining.

Table 2. Main metrics and interpretation.

Metric	Direction	Definition	Caveat
Productivity	↑	Share of submitted bills enacted.	Can reward bad throughput.
Yield	↑	Generated public benefit per submitted proposal.	Synthetic and generator-dependent.
Revision moderation	↑	Score combining moderation, legitimacy, and incrementalism.	Not a neutral welfare measure.
Public support	↑	Enacted-bill support composite from generated support signals.	Not observed public opinion.
Low-support enactment	↓	Enacted bills below 50% generated support.	Depends on synthetic support threshold.
Risk control	↑	Inverse composite of harm, capture, low support, proposer gain, and policy shock.	Composite score.
Admin feasibility	↑	1 – procedural-load index from attention, review, challenge, amendment, and selection costs.	Indirect cost proxy.

4 Metrics and Campaign

The simulator reports a dashboard rather than a single welfare score. Table 2 defines the metrics used in the main paper; secondary diagnostics remain in generated reports and appendix tables.

Yield prevents a system that blocks nearly all proposals from looking equivalent to one that produces generated public value. Risk control averages inverse low-support passage, minority harm, lobbying capture, public-preference distortion, concentrated-harm passage, proposer gain, and policy shock terms; administrative feasibility is 1 – the procedural-load index.

As one sensitivity check, the paper reports a revision-moderation-weighted value profile combining revision moderation, productivity, risk control, administrative feasibility, and public support. The main interpretation relies on metric-level tradeoffs rather than this scalar score.

All main results come from one reported comparison campaign. It combines broad assumptions, adversarial generator cases, weighted party-system cases, and a stylized rising-contention timeline in one CSV. The canonical run uses 120 generated worlds per case, 101 legislators, 60 bills per run, and seed 20260428. Cases vary polarization, party loyalty,




Fig. 1. Modular simulation pipeline. Each institutional bundle intervenes at one or more stages while generated legislatures, bill streams, and status-quo-relative voting remain shared across comparisons.

revision-moderation culture, lobbying, constituency pressure, party count, proposal pressure, capture pressure, and anti-lobbying reform demand. The adversarial bill-stream cases stress extremity, capture, delayed benefits, concentrated harm, and backlash dynamics.

Scenario selection follows a fixed coverage rule: include the stylized U.S.-like benchmark, simple majority, committee regular order, discharge petition, open floor calendar, a combined content-selection family, anti-capture access, harm-weighted majority, one burden-shifting stress case, one burden-shifting mediation case, and a portfolio hybrid. The appendix reports the broader scenario catalog and fixed-rule family champions.

Table 3. Mechanism families used in the main comparison. Labels match the Results figures; the appendix reports the broader implemented catalog.

Family	Labels	Mechanism theme	Primary stress risk
Conventional benchmark	CUR	House/Senate/veto gate	Bottlenecks
Simple majority	SM	Affirmative threshold	Low screening
Committee regular order	COMM	Committee-first gate	Burial/capture
Discharge petition	DIS	Committee bypass	Bypass pressure
Open floor calendar	OPEN	Default-open agenda	Floor overload
Content selection	SEL	Alternatives/tournament	Clones/agenda manipulation
Anti-capture access	ACG	Pre-floor capture screen	Gate bias
Harm-weighted majority	HARM	Affected-group threshold	Holdouts
Burden-shifting stress	DP	Pass unless blocked	False silence
Burden shift + mediation	DPM	Mediation plus challenge	Cost/weak consent
Portfolio hybrid	PORT	Combined safeguards	Complexity

The chamber-structure supplement tests representation architecture, committee power, eligibility filters, and ex ante review. It is separate because it asks how those structures shift the same productivity/revision-moderation tradeoff.

Empirical boundary checks are artifact screens, not model validation in the statistical or causal sense. make calibrate compares conventional baselines with broad benchmark ranges derived from Voteview, Congress.gov, govinfo, Comparative Agendas Project, lobbying-disclosure and OpenFEC records, Center for Effective Lawmaking measures, SCDB merits cases, Federal Register final-rule records, Congress.gov public-law text flags, and institutional indicators [Comparative Agendas Project 2026; Federal Election Commission 2026; Lewis et al. 2021; Library of Congress 2026; United States Government Publishing Office 2020; United States Senate Office of Public Records 2026; V-Dem Institute 2026; Volden and Wiseman 2014]. The legislative-flow layer applies one reviewed classifier to every direct action in complete pinned 116th-, 117th-, and 118th-Congress govinfo H.R. and S. BILLSTATUS archives: respectively 14,148/68,345, 15,066/72,047, and 16,213/75,239 bills/actions. A deterministic 117th-Congress bill-ID split supplies 7,564 calibration bills and 7,502 held-out bills. Only the calibration half selects the current-Congress workflow’s calendar-priority threshold from a fixed 0.60–0.74 grid and 50-seed panel; floor consideration and enactment enter selection, while committee advancement is an upstream check. The selected 0.68 threshold yields committee-advance, substantive-floor, and enactment rates of 0.106, 0.061, and 0.027, compared with held-out rates of 0.101, 0.068, and 0.025; all 50 leave-one-seed-out panels reselect it. Without refitting, all three rates pass their frozen 116th-Congress backcast tolerances; the 118th forecast passes committee and floor tolerances but overpredicts enactment by 0.010825, missing its 0.010 tolerance by 0.000825. Thus 5 of 6 external cohort-metric cells pass. Classifier v3 adds successful override evidence without changing established funnel counts.

A separate mechanism test joins the complete 116th–118th-Congress H.R. panels to the official House Rules Committee activity surveys [U.S. House of Representatives, Committee on Rules 2021, 2023, 2025]. The source study retains 29,335 bills, 458 literal Table 1a rule grants, and four mutually exclusive observed floor routes. The plan’s asserted pre-fit timing is not independently established in the repository. A post-source-audit plan fixes 1,445 threshold triples, uses the 116th–117th Congresses for development, and reserves fresh simulation seeds for an unchanged 118th-Congress test. The selected calendar, special-rule, and suspension thresholds are 0.680, 0.475, and 0.750; 16 of 20 leave-one-seed-out panels reselect that triple. Committee advancement, floor consideration, and restrictive special-rule share pass their test tolerances, but route total variation is 0.110470 against the locked 0.100 tolerance. The suspension threshold is at the grid boundary, and the simulator underproduces special-only routes and overproduces mixed and other routes. We retain the 3-of-4 gate failure without expanding the grid or retuning on the test cohort. This falsification result indicates an incomplete route mechanism and rules out interpreting the benchmark as a calibrated model of House agenda control.

Separate source-pinned panels retain 4,021 H.R./S. and 187 H.J.Res./S.J.Res. presentments across the 108th–118th Congresses. Their exact 47-veto and six-override set matches official Senate histories, while president and chamber-control fields follow the official House history [Office of the Historian, U.S. House of Representatives 2026; United States Senate 2026]. The combined empirical veto rate is 0.011169 versus 647 vetoes among 2,621 frozen simulator decisions, or 0.246852 and an exact-count ratio of 22.101. A final-vote panel preserves one House and one Senate approval row per measure, including 1,685 official final roll calls and 6,731 explicit nonrecorded approvals. After source construction and descriptive audit but before fitting, a committed plan fixed a low-event L2-penalized logistic model, explicit roll-call-recording indicators, no support imputation, proper probability scores [Gelman et al. 2008; Gneiting and Raftery 2007], and whole-Congress temporal tests. The primary model trains on 3,921 measures through the 117th Congress and tests 287 118th-Congress measures with 13 vetoes. The structural M1 ablation reaches log loss 0.058347

Table 4. Selected empirical flow smoke tests. Raw inputs do not validate the synthetic welfare, revision moderation, harm, capture, or public-support metrics.

Signal	Raw	Simulator proxy	Does not validate
Bill attrition	enactmentRate=0.024	current-congress-workflow/productivity	public benefit or law quality
Floor consideration	floorLoad=0.067	current-congress-workflow/floor	agenda fairness or policy merit
Committee advancement from bill actions	committeeAdvanceRate=0.100	current-congress-workflow/committeeAdvanceRate	committee expertise or capture
Committee reporting from bill actions	committeeReportRate=0.064	—	committee expertise or capture
Roll-call coalition size	coalitionSize=0.604	current-system/averageEnactedSupport	public opinion
Sponsor proposal concentration	sponsorIntroductionGini=0.389	current-system/proposerAccessGini	capture or representational equity

and Brier score 0.015232. M2 reaches 0.026977 and 0.008143 versus 0.223369 and 0.044586 for training prevalence; calibration-in-the-large is 0.003149, so the frozen two-part gate passes. In the secondary 116th-Congress check, M1 and M2 log losses are 0.062969 and 0.038738 versus 0.145942, with M2 calibration-in-the-large 0.000908. This is narrow predictive transport, not a blinded preregistration or a causal model: a post-fit audit finds 12 of 13 primary test vetoes among only 17 joint resolutions. The stress simulator remains badly mis-scaled for U.S. veto frequency, and transport beyond this concentrated measure mix is untested.

Separate held-out and flow checks use documented empirical comparison samples for sponsors, roll calls, lobbying, campaign finance, district opinion, judicial review, implementation, statutory revision, and comparative institutions. The district-opinion path begins with a bounded Cumulative CES aggregate [Kuriwaki 2025], sponsor-district bill metadata, policy-area context, Census/ACS district context, and a 22-bill review queue. An official bill-text review retains one directionally related historical policy-preference item for H.R. 8404 and preserves 21 negative dispositions against forced policy-area matches. That retained item is joined to pinned 2012 and 2016 CCES Common Content geography and validated-voter weights [Ansolahehere and Schaffner 2012, 2016; Dagonel 2023], producing two privacy-thresholded NY-10 direct-weighted aggregates with 372 published responses. The item does not use the bill wording, both estimates predate enactment, and the layer supplies neither design-based uncertainty, MRP, affected-group evidence, causal representation, nor model validation. Across all source families, 14 targeted held-out checks pass broad prespecified ranges; the linkage audit reports all fourteen source families as linked, metadata-linked, or partially linked. Passing these screens is minimum plausibility evidence only.

A post-fit House source audit finds one Senate action misattributed as House suspension evidence. A direct-action-only screen increases fixed-model route variation to 0.111950; the original fit and failed gate remain unchanged. Retained House Clerk and resolution records establish a separate indirect suspension path through H. Res. 1061 for that bill [Office of the Clerk, U.S. House of Representatives 2024; U.S. House of Representatives 2024]. The screen is not a complete route correction. History-wide source routes and the model’s single-decision score overlap are not equivalent procedural sequences.

The fitted special-rule and suspension thresholds assign diagnostic labels to admitted bills; they do not alter the inner voting threshold or amendment rights. Matching these labels would not establish implementation of actual House procedures.

5 Results

The selected mechanism-family comparison shows three patterns. First, throughput and support/risk diagnostics separate: mechanisms that enact more proposals can also enact more low-public-support proposals. Second, content-selection stages change the collective intelligence process by altering which alternatives enter final aggregation, not merely by changing the final yes/no threshold. Because pairwise alternatives and amendment tournaments produce

nearly identical aggregate profiles in this campaign, the main text reports them as a single content-selection family. Third, portfolio safeguards reduce several risk diagnostics but increase procedural load. The stylized U.S.-like benchmark has low productivity but stronger revision moderation, risk, and support diagnostics than its directional profile suggests.

The empirical boundary checks are narrow screens for observable legislative-flow and roll-call proxies. For the richer current-Congress workflow, the full bill census now disciplines committee advancement, floor consideration, and enactment simultaneously; it does not validate generated public benefit, revision moderation, harm, capture, representation, or public-support diagnostics.

The main robustness table reports numeric generator-sensitivity, vote-kernel-ablation, and seed-stability checks. Detailed Pareto checks, ablation deltas, manipulation-stress losses, and separate PAIR/AMT rows appear in the appendix. The generator-sensitivity rows and fairness controls show that the framework can expose dependence on generator assumptions, but they do not remove that dependence. These results are mechanism-behavior hypotheses under the implemented generator, not general rankings of legislative institutions.

The chamber-structure supplement is reported in the appendix. It asks a separate question about representation architecture and review design, so it is not part of the ACM paper’s central evidence chain.




Fig. 2. Productivity tradeoff views for the representative systems. Panel (a) plots risk control across the main broad/adversarial campaign averages, with thin five-seed uncertainty bars where visible; panel (b) plots revision moderation under weighted party-system sensitivity cases. Higher is better on both y-axes. Labels are the mechanism-family labels from Table 3. The asterisk and larger square mark the stylized U.S.-like benchmark.

The burden-shifting stress cases illustrate the same separation. Open burden-shifting passage has high productivity but weak revision-moderation and support diagnostics, while mediation improves some content metrics at additional procedural cost. These results are best read as failure-mode probes rather than as a central reform comparison.

Manipulation-stress cases use bounded adversaries rather than fully co-evolving strategic parties. The adversary can choose clone or decoy proposals, noisier citizen-panel information, loose harm claims, astroturf objection pressure, proposal flooding, or anti-lobbying backlash under fixed information and budget assumptions. Stress objectives are enactment of low-support bills, blockage of high-benefit bills, or administrative overload. Appendix tables report paired losses, while adaptive worst-case search remains future work.

The portfolio variants illustrate the tradeoff between combining safeguards and increasing administrative load. They improve several risk diagnostics but do not dominate simpler alternative-selection mechanisms.

Table 5. Selected metric profiles for the main mechanisms.

Label	Prod.	Rev. mod.	Public	Risk	Admin feas.	Example profile
CUR*	0.05	0.36	0.77	0.93	0.98	0.57
SM	0.33	0.24	0.68	0.85	0.93	0.55
COMM	0.21	0.28	0.73	0.89	0.97	0.56
DIS	0.14	0.31	0.75	0.92	0.92	0.56
OPEN	0.36	0.25	0.67	0.84	0.90	0.55
SEL	0.55	0.50	0.79	0.94	0.71	0.67
ACG	0.27	0.24	0.70	0.88	0.94	0.55
HARM	0.28	0.26	0.69	0.87	0.93	0.55
DP	0.87	0.14	0.57	0.63	0.93	0.57
DPM	0.81	0.24	0.62	0.73	0.79	0.59
PORT	0.53	0.43	0.76	0.93	0.76	0.65

Note: Values are normalized to the metric direction shown in Table 2; grayscale shading is secondary. Admin feas. is administrative feasibility, the inverse of the cost index. Example profile is 0.30Rev. mod. +0.20Prod. +0.20Risk +0.15Admin feas. +0.15Public. The asterisk marks the stylized U.S.-like benchmark.

Table 6. Main robustness checks with numeric evidence.

Check	SEL	PORT	DP/SM	Rank ρ	Main implication
High-benefit extreme	0.69	0.72	0.57	0.95	Moderation-friendly generator is relaxed by allowing high-distance bills to carry high generated value.
Revision dilution	0.67	0.71	0.67	0.78	Moving toward the center can reduce generated public benefit.
Lobby information	0.74	0.71	0.28	0.70	Organized lobbying can supply useful technical information rather than only capture pressure.
Public-opinion error	0.68	0.68	0.41	0.93	Generated support and generated benefit are decoupled by noisy or biased support signals.
Minority-rights harm	0.66	0.64	0.19	0.57	Majority-supported bills can impose concentrated rights-like harm.
No support/revision vote	0.57	—	SM 0.49	—	Vote-kernel ablation removes public-support and revision-moderation vote inputs while retaining them as diagnostics.
Five-seed sweep	0.699–0.701	0.700–0.701	0.733–0.734	n=5	Entries show directional-score min–max across independent base seeds; ranges are small relative to mechanism-family differences.

Note: SEL and PORT entries report the example profile except in the seed row, where they report directional-score min–max. DP/SM reports default-pass low-public-support enactment for generator checks, the no-signal simple-majority profile for the vote-kernel ablation, and default-pass seed range for the seed sweep. Rank ρ is the Spearman correlation of displayed mechanism-family profile ranks against the main broad/adversarial average.

The reported ordering should therefore be read as a hypothesis about mechanism behavior under the implemented generator, not as an institutional ranking.

6 Limitations

The model is stylized. It uses one primary policy dimension, generated legislators, generated support and benefit, simplified issue domains, and abstract lobbying units. The multidimensional package scenario is a cross-issue bargaining proxy, not a full two-dimensional ideal-point model; claims about alternatives, tournaments, revision moderation, and package bargaining are therefore intentionally narrow. Public-support diagnostics distinguish salience and costly affected-group pressure from cheap amplified volume, but they are still proxies. Because public benefit is internal to the generator, welfare comparisons are conditional on model assumptions. Systems that screen extreme, uncertain, or

lobby-heavy bills may look better partly because the generator often encodes those features as lower expected public value. Adversarial cases partially relax this assumption, but welfare remains a synthetic signal.

The empirical layer screens conventional baselines against empirical ranges, and the held-out and flow smoke tests cover documented empirical inputs. That is enough to make several agenda-access, legislative-flow, roll-call, campaign-finance, district-opinion, review, implementation-delay, revision-text, and comparative-profile proxies inspectable, but it is not enough for model validation or parameter fitting. The linkage report currently documents complete govinfo BILLSTATUS lifecycle rows for all 14,148 116th-, 15,066 117th-, and 16,213 118th-Congress H.R. and S. bills, plus independent Congress.gov-to-govinfo agreement for 180 of 180 bounded 118th-Congress rows; bounded LDA issue context for 144 of 146 lobbying rows and 484 exact activity-text bill mentions across 26 cached public-law bill IDs; an FEC recipient-metadata join for 192 of 194 campaign-finance rows plus bounded issue, member, district, and sponsored-bill context; and sponsor-district public-law metadata for 63 of 1305 district-opinion row keys. The public-opinion review covers 66 policy-context rows across 22 bills, source-reviews official bill text for all 22, retains one historical related-issue item, preserves 21 negative dispositions, and publishes two annual privacy-thresholded aggregates for the retained item. Twenty-one bills remain proxy-only, exact or contemporaneous bill support remains absent, and affected-group support/harm rows remain at zero. Official Census TIGERweb denominators and broad ACS context cover all 22 packets across 21 sponsor districts. The report also records SCDB/Federal Register U.S.C.-section overlaps for 26 of 9341 merits rows, Federal Register document/docket metadata for 495 of 500 implementation rows, authority-text matches for 15 of 40 cached public-law rows, proposed-rule history for 23 of 51 authority-matched final-rule rows, and Congress.gov bill/action metadata for 40 of 120 public-law revision rows. The bill-law spine attaches same-policy campaign-finance sponsor context to 17 public-law rows, same-policy LDA issue/bill context to 36 public-law rows, exact official LDA filing-text bill identifiers for 26 public-law rows, exact LDA bill/action metadata context for those 26 rows, and bounded stored LDA activity-text signals; these are context, identifier, text-signal, activity-text disposition, and metadata evidence, not bill-specific influence. Separate govinfo, court-law, rulemaking-authority, rulemaking-history, bill-law, district-public-opinion policy-context, bill-item alignment, historical bill-topic support, Census denominator, ACS context, survey-source crosswalk, lobbying, campaign-finance, and bill-finance/lobbying reports expose these bounded joins and the missing evidence that remains. Exact or contemporaneous bill-topic district opinion beyond the historical related-issue pilot, bill-text-specific affected-population detail, affected-group representation, lobbying sponsor/member targeting and influence, bill-linked campaign finance, source-specific agenda evidence, presidential-choice transport beyond the concentrated 118th-Congress pattern, direct case-to-public-law review, emergency court review, fuller implementation feedback, full statutory lineage, and observed cross-national institutional productivity or fit remain validation gaps; findings that depend on those components should be read as synthetic design hypotheses. The Cumulative CES district aggregate remains a direct survey proxy. The two new annual estimates add historical related-issue district context only; they are not MRP, exact or contemporaneous bill support, affected-group harm evidence, causal representation, or model validation. The older committee and sponsor summaries are derived from sampled Congress.gov bill actions and remain rough. The three govinfo censuses replace the sample as the central legislative-flow benchmark, but strict committee reporting remains descriptive because the simulator maps only the broader ordered-reported, reported, or discharged advancement union. The OpenFEC sample supports concentration, outside-spending summaries, public recipient metadata, bounded issue-sector context, bounded candidate-member context, bounded House-candidate district context, bounded candidate-to-sponsored-bill context, public target-scope review, and cached committee/floor bill-action flags for queued rows only, not causal influence, committee-action influence, legislative-outcome causality, or bill-level capture. The LDA sample supports

spending concentration, issue-taxonomy context, policy-area bill/action context, exact filing-text bill identifiers, exact bill/action metadata context, bounded stored activity-text support/position signals, a manual-review queue, a 4-row high-priority manual review slice, 102 medium-priority review packets, source-reviewed directional and position/activity packet dispositions, external current-bill mention-packet review, and cached committee/floor bill-action flags for a public-law subset only, not lobbying-contact confirmation, sponsor targeting beyond activity-text references, committee-action influence, roll-call influence, legislative-outcome causality, public benefit, or capture validation. Congress.gov public-law revision rows support only amendment, reauthorization/extension, repeal, sunset/expiration text flags, bounded bill/action metadata, same-policy finance/lobbying context, exact LDA bill identifiers and bill/action context, authority-search metadata, proposed-history/comment-portal metadata, timing metadata, bounded court U.S.C.-section overlap metadata, and the derived spine inventory, not codified lineage, complete comments, implementation outcomes, bill-specific finance/lobbying influence, or direct court invalidation. Fixture samples exercise adapter shapes but are not validation data. Strategic behavior remains bounded: proposers and lobby groups adapt, and norm erosion can raise delay, but parties, committees, coalitions, courts, agencies, media, and elections do not yet co-evolve. Chamber scenarios are sensitivity prototypes. Citizen panels, initiatives, objection windows, tournaments, audits, challenge tokens, judicial review, and law-registry review are optimistic prototypes whose costs are represented only indirectly. Package bargaining models side payments, delay, jurisdictional trades, and distributive load, not legal text or appropriations.

7 Reproducibility

The submitted manuscript and supplement are anonymized. The Java 21 artifact uses Makefile entrypoints and tracked or generated CSV summaries. `make reproduce-paper-offline` is the canonical no-network target: from fixed seeds and committed empirical panels it regenerates the campaign, diagnostics, figures, both PDFs, and all publication checks. `make supplement-anonymous` builds the supplement, and `make test` runs the simulator and validation unit suite.

The empirical chain has three narrower gates. `make govinfo-bill-census-check` verifies pinned lifecycle and presidential-decision panels, action hashes, classifier cases, final-vote arithmetic, and frozen transport results. `make presidential-choice-study-check` rebuilds the locked temporal model and verifies source, specification, prediction, metric, and report hashes. `make validation-gap-report` regenerates source-family inventories, linkage boundaries, and the appendix table. The README and Makefile list source-specific refresh and review targets. Network-dependent refreshes remain optional and outside the no-network path; fixture samples never substitute for raw empirical inputs. Reproduced values must match the tracked CSV and PDF manifests.

Use of Large Language Model Tools

The research question and core simulator idea originated with the human author. OpenAI ChatGPT/Codex tools available in the authoring environment during the April–May 2026 revision period were used for grammar and style editing, restructuring suggestions, code debugging assistance, and drafting assistance for portions of the simulator implementation, generated-report scripts, and manuscript prose. LLM tools were not used as authors. Tool outputs were edited and checked by the author; references, code outputs, experiments, tables, figures, and claims were manually reviewed against the repository artifacts before inclusion. No large language model is listed as an author.

8 Conclusion

This paper contributes a reproducible simulation framework for comparing legislative mechanism bundles as collective intelligence architectures. The reusable contribution is an experimental architecture for making legislative collective

decision tradeoffs inspectable under shared synthetic worlds. The results support the framework’s main methodological claim: legislative productivity, revision moderation, public support, risk control, and administrative feasibility separate under different institutional mechanisms. In the reported synthetic campaigns, content-selection stages often change the productivity/risk surface more than threshold-only baselines, while portfolio variants illustrate the cost of combining safeguards. Future work should replace synthetic proxies with stronger public-opinion, bill-linked lobbying and campaign-finance, source-specific agenda data, presidential-choice replication beyond the current concentrated temporal test, direct court-review and emergency-order, implementation, full statutory-lineage beyond bill/action metadata, observed comparative-productivity, and bargaining data.

References

- Stephen Ansolabehere and Brian Schaffner. 2012. CCES Common Content, 2012. doi:10.7910/DVN/HQEVPK
- Stephen Ansolabehere and Brian F. Schaffner. 2016. CCES Common Content, 2016. doi:10.7910/DVN/GDF6Z0
- David Austen-Smith. 1993. Information and Influence: Lobbying for Agendas and Votes. *American Journal of Political Science* 37, 3 (1993), 799–833. doi:10.2307/2111575
- David Austen-Smith. 1995. Campaign Contributions and Access. *American Political Science Review* 89, 3 (1995), 566–581. doi:10.2307/2082974
- David P. Baron and John A. Ferejohn. 1989. Bargaining in Legislatures. *American Political Science Review* 83, 4 (1989), 1181–1206. doi:10.2307/1961664
- Alessandra Casella. 2005. Storable Votes. *Games and Economic Behavior* 51, 2 (2005), 391–419.
- Comparative Agendas Project. 2026. Datasets and Codebooks. https://www.comparativeagendas.net/datasets_codebooks
- Gary W. Cox and Matthew D. McCubbins. 2005. *Setting the Agenda: Responsible Party Government in the U.S. House of Representatives*. Cambridge University Press, Cambridge, United Kingdom.
- Angelo Dagonel. 2023. Cumulative CES Policy Preferences. doi:10.7910/DVN/OSXDQO
- John M. de Figueiredo and Brian Kelleher Richter. 2014. Advancing the Empirical Research on Lobbying. *Annual Review of Political Science* 17 (2014), 163–185. doi:10.1146/annurev-polisci-100711-135308
- Scott de Marchi and Scott E. Page. 2014. Agent-Based Models. *Annual Review of Political Science* 17 (2014), 1–20. doi:10.1146/annurev-polisci-080812-191558
- Anthony Downs. 1957. *An Economic Theory of Democracy*. Harper and Row, New York, NY.
- Joshua M. Epstein. 2006. *Generative Social Science: Studies in Agent-Based Computational Modeling*. Princeton University Press, Princeton, NJ.
- Federal Election Commission. 2026. OpenFEC API Documentation. <https://api.open.fec.gov/developers/>
- James S. Fishkin. 2009. *When the People Speak: Deliberative Democracy and Public Consultation*. Oxford University Press, Oxford, United Kingdom.
- Andrew Gelman, Aleks Jakulin, Maria Grazia Pittau, and Yu-Sung Su. 2008. A Weakly Informative Default Prior Distribution for Logistic and Other Regression Models. *The Annals of Applied Statistics* 2, 4 (2008), 1360–1383. doi:10.1214/08-AOAS191
- Jacob E. Gersen. 2007. Temporary Legislation. *University of Chicago Law Review* 74, 1 (2007), 247–298.
- Nigel Gilbert and Klaus G. Troitzsch. 2005. *Simulation for the Social Scientist* (second ed.). Open University Press, Maidenhead, United Kingdom.
- Martin Gilens and Benjamin I. Page. 2014. Testing Theories of American Politics: Elites, Interest Groups, and Average Citizens. *Perspectives on Politics* 12, 3 (2014), 564–581. doi:10.1017/S1537592714001595
- Thomas W. Gilligan and Keith Krehbiel. 1987. Collective Decisionmaking and Standing Committees: An Informational Rationale for Restrictive Amendment Procedures. *Journal of Law, Economics, and Organization* 3, 2 (1987), 287–335.
- Tilman Gneiting and Adrian E. Raftery. 2007. Strictly Proper Scoring Rules, Prediction, and Estimation. *J. Amer. Statist. Assoc.* 102, 477 (2007), 359–378. doi:10.1198/016214506000001437
- Volker Grimm, Uta Berger, Donald L. DeAngelis, J. Gary Polhill, Jarl Giske, and Steven F. Railsback. 2010. The ODD Protocol: A Review and First Update. *Ecological Modelling* 221, 23 (2010), 2760–2768. doi:10.1016/j.ecolmodel.2010.08.019
- Richard L. Hall and Alan V. Deardorff. 2006. Lobbying as Legislative Subsidy. *American Political Science Review* 100, 1 (2006), 69–84. doi:10.1017/S0003055406062010
- Richard L. Hall and Frank W. Wayman. 1990. Buying Time: Moneyed Interests and the Mobilization of Bias in Congressional Committees. *American Political Science Review* 84, 3 (1990), 797–820. doi:10.2307/1962767
- Keith Krehbiel. 1991. *Information and Legislative Organization*. University of Michigan Press, Ann Arbor, MI.
- Keith Krehbiel. 1998. *Pivotal Politics: A Theory of U.S. Lawmaking*. University of Chicago Press, Chicago, IL.
- Shiro Kuriwaki. 2025. Cumulative CES Common Content. doi:10.7910/DVN/II2DB6
- Steven P. Lalley and E. Glen Weyl. 2018. Quadratic Voting: How Mechanism Design Can Radicalize Democracy. *AEA Papers and Proceedings* 108 (2018), 33–37. doi:10.1257/pandp.20181002
- Jeffrey B. Lewis, Keith Poole, Howard Rosenthal, Adam Boche, Aaron Rudkin, and Luke Sonnet. 2021. Voteview: Congressional Roll-Call Votes Database. <https://voteview.com/>
- Library of Congress. 2026. Congress.gov API and Legislative Data. <https://www.congress.gov/help/using-data-offsite>

- Richard D. McKelvey. 1976. Intransitivities in Multidimensional Voting Models and Some Implications for Agenda Control. *Journal of Economic Theory* 12, 3 (1976), 472–482. doi:10.1016/0022-0531(76)90040-5
- Birgit Müller, Friedrich Bohn, Gunnar Dreßler, Jürgen Groeneveld, Christian Klassert, Romina Martin, Martin Schlüter, Johannes Schulze, Helge Weise, and Nils Schwarz. 2013. Describing Human Decisions in Agent-Based Models – ODD + D, an Extension of the ODD Protocol. *Environmental Modelling and Software* 48 (2013), 37–48. doi:10.1016/j.envsoft.2013.06.003
- Office of the Clerk, U.S. House of Representatives. 2024. Roll Call 64: H. Res. 1061, Motion to Suspend the Rules and Agree. <https://clerk.house.gov/Votes/202464>
- Office of the Historian, U.S. House of Representatives. 2026. Party Government Since 1857. <https://history.house.gov/Institution/Presidents-Coinciding/Party-Government/>
- Organisation for Economic Co-operation and Development. 2020. *Innovative Citizen Participation and New Democratic Institutions: Catching the Deliberative Wave*. OECD Publishing, Paris, France. doi:10.1787/339306da-en
- Sofia Ranchordas. 2014. *Constitutional Sunsets and Experimental Legislation: A Comparative Perspective*. Edward Elgar Publishing, Cheltenham, United Kingdom.
- Thomas Romer and Howard Rosenthal. 1978. Political Resource Allocation, Controlled Agendas, and the Status Quo. *Public Choice* 33, 4 (1978), 27–43. doi:10.1007/BF03187594
- Kenneth A. Shepsle. 1979. Institutional Arrangements and Equilibrium in Multidimensional Voting Models. *American Journal of Political Science* 23, 1 (1979), 27–60. doi:10.2307/2110770
- Kenneth A. Shepsle and Barry R. Weingast. 1987. The Institutional Foundations of Committee Power. *American Political Science Review* 81, 1 (1987), 85–104. doi:10.2307/1960780
- T. N. Tideman. 1987. Independence of Clones as a Criterion for Voting Rules. *Social Choice and Welfare* 4, 3 (1987), 185–206. doi:10.1007/BF00433944
- George Tsebelis. 2002. *Veto Players: How Political Institutions Work*. Princeton University Press, Princeton, NJ.
- United States Government Publishing Office. 2020. Bill Status XML Bulk Data. <https://www.govinfo.gov/features/bill-status-xml-bulk-data>
- United States Senate. 2026. Summary of Bills Vetoes, 1789-Present. <https://www.senate.gov/legislative/vetoes/vetoCounts.htm>
- United States Senate Office of Public Records. 2026. Lobbying Disclosure Act Database. <https://lda.senate.gov/system/public/>
- U.S. House of Representatives. 2024. H. Res. 1061, Engrossed in House, March 6, 2024. <https://www.govinfo.gov/content/pkg/BILLS-118hres1061eh/html/BILLS-118hres1061eh.htm>
- U.S. House of Representatives, Committee on Rules. 2021. *Survey of Activities of the House Committee on Rules for the 116th Congress*. Technical Report House Report 116-722. U.S. Government Publishing Office. <https://www.govinfo.gov/content/pkg/CRPT-116hrpt722/html/CRPT-116hrpt722.htm>
- U.S. House of Representatives, Committee on Rules. 2023. *Survey of Activities of the House Committee on Rules for the 117th Congress*. Technical Report House Report 117-709. U.S. Government Publishing Office. <https://www.govinfo.gov/content/pkg/CRPT-117hrpt709/html/CRPT-117hrpt709.htm>
- U.S. House of Representatives, Committee on Rules. 2025. *Survey of Activities of the House Committee on Rules for the 118th Congress*. Technical Report House Report 118-979. U.S. Government Publishing Office. <https://www.govinfo.gov/content/pkg/CRPT-118hrpt979/html/CRPT-118hrpt979.htm>
- V-Dem Institute. 2026. V-Dem Dataset. <https://www.v-dem.net/data/the-v-dem-dataset/>
- Craig Volden and Alan E. Wiseman. 2014. *Legislative Effectiveness in the United States Congress: The Lawmakers*. Cambridge University Press, New York, NY.
- Barry R. Weingast and William J. Marshall. 1988. The Industrial Organization of Congress; or, Why Legislatures, Like Firms, Are Not Organized as Markets. *Journal of Political Economy* 96, 1 (1988), 132–163. doi:10.1086/261528